



Languages for concurrency and
distribution

Concurrency in CUDA



Table of content

- ❑ What is CUDA?**
- ❑ Heterogeneous Computing**
- ❑ CUDA Paradigm**
- ❑ Threads Communication**
- ❑ Dynamic Parallelism**
- ❑ Final consideration**





What is CUDA?

CUDA (Compute Unified Device Architecture)



- Parallel computing model
- GPU based
- Developed by Nvidia (proprietary)
- Available for many languages (CUDA C/C++, pyCUDA, CUDA Fortran, CUDA Matlab...)

CUDA c/c++

- Based on C/C++ standard
- NVCC compiler
- Small set of extensions to allow heterogeneous computing




Heterogeneous Computing

CPU vs GPU

CPUs: Latency Oriented

- Powerful ALU
- Large caches
- Sophisticated control:
 - Branch prediction
 - Data forwarding

Faster on sequential code




Host: CPU and its memory(host memory)

● ALU
● Control
● Cache
● DRAM

GPUs : Throughput Oriented


- Small caches
- Simple control
- Energy efficient ALU
- Require massive number of threads to tolerate latency

Faster on parallel code



Device: GPU and its memory(device memory)

Processing workFlow



- 1) Copy input data from CPU to GPU memory
- 2) Load GPU code and execute it, caching data on chip for performance
- 3) Copy results from GPU memory to CPU memory

Device Code

Host Code

```
#include <iostream>
#include <cuda.h>

// CUDA kernel (Device code) that runs on the GPU
__global__ void square(float *d_array, int size) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < size) {
        d_array[idx] = d_array[idx] * d_array[idx];
    }
}

int main() {
    /*-----1-----*/
    const int arraySize = 10; // Array size
    float *d_array, *h_array;
    h_array = (float*)malloc(arraySize*sizeof(float)); // Host array

    // Allocate device memory
    cudaMalloc(&d_array, arraySize * sizeof(float)); // Device array

    // Initialize array on host
    for (int i = 0; i < arraySize; i++) {
        h_array[i] = i + 1; // Array of {1, 2, 3, ..., 10}
    }

    //Copy the host data into device memory
    cudaMemcpy(d_array, h_array, arraySize * sizeof(float), cudaMemcpyHostToDevice);

    /*-----2-----*/
    //Launch kernel on the device with 1 block and 10 threads(one per array element)
    square<<<1, arraySize>>>(d_array, arraySize);

    /*-----3-----*/
    // Copy the result back to the host
    cudaMemcpy(h_array, d_array, arraySize * sizeof(float), cudaMemcpyDeviceToHost);
    cudaDeviceSynchronize();

    // Free device memory
    cudaFree(d_array);
    free(h_array);
}
```


In the code

- **CudaMalloc()** : Allocate memory on the device.
- **CudaMemcpy()** : copy values from host to device(and viceversa).
- **CudaFree()** : Frees memory that was previously allocated on the device.



CUDA Paradigm



Threads Hierarchy



Kernel

- Functions or full programs.
- Executed in parallel.
- Define hierarchy of grids of thread blocks.
- Support a synchronization mechanism.
- Threads are identified by their thread index (*threadIdx*), block index (*blockIdx*).

```
global void kernelEx(params){...}  
kernelEx<<<(3,2),(2,2)>>>(params);
```



Kernel in the code

- Keyword **`__global__`** in the function header.
- `kName<<<gridDim, blockDim>>>(params)` to launch the kernel from the CPU.
- **`cudaDeviceSynchronize()`** blocks the host until all previously launched CUDA tasks on the device have completed execution

Example : array addition

```
#include <iostream>

#include <cuda_runtime.h>

#define N 1024 // Total number of elements in the array
#define THREADS_PER_BLOCK 256 // Number of threads per
block
#define BLOCKS 4 // Number of threads per block

// DEVICE code
// CUDA Kernel for vector addition
__global__ void add(int *a, int *b, int *c, int n) {
    // Calculate the global index for each thread
    int index = threadIdx.x + blockIdx.x * blockDim.x;

    // Ensure that we don't go out of bounds
    if (index < n) {
        c[index] = a[index] + b[index];
    }
}
```

```
// HOST code
int main() {
    // Host arrays
    int h_a[N], h_b[N], h_c[N];

    // Initialize the host arrays with some values
    for (int i = 0; i < N; i++) {
        h_a[i] = i;
        h_b[i] = i * 2;
    }

    // Device arrays
    int *d_a, *d_b, *d_c;

    // Allocate memory on the GPU for the device arrays
    cudaMalloc((void **)&d_a, N * sizeof(int));
    cudaMalloc((void **)&d_b, N * sizeof(int));
    cudaMalloc((void **)&d_c, N * sizeof(int));

    // Copy data from the host arrays to the device arrays
    cudaMemcpy(d_a, h_a, N * sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, h_b, N * sizeof(int), cudaMemcpyHostToDevice);

    // Launch the kernel with multiple blocks and threads
    add<<<BLOCKS, THREADS_PER_BLOCK>>>(d_a, d_b, d_c, N);

    // Wait for the GPU to finish executing the kernel
    cudaDeviceSynchronize();

    // Copy the result from the device array d_c to the host array h_c
    cudaMemcpy(h_c, d_c, N * sizeof(int), cudaMemcpyDeviceToHost);

    // Print the result
    std::cout << "Result of vector addition:\n";
    for (int i = 0; i < N; i++) {
        std::cout << h_a[i] << " + " << h_b[i] << " = " << h_c[i] << "\n";
    }

    // Free the device memory
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    return 0;
}
```





Threads communication

Communication and memory

- *Shared Memory*: small amount of data
 - Only from threads of the same block
 - Use Streaming Multiprocessor(SM)
 - Space splitted equally among threadsblocks → limited amount of memory
- *Global Memory*: for bigger amount data
 - Any thread
 - Slower
 - Synchronization can be an issue

Kernel Structure is Crucial!

Communication in the code

- **`__shared__`** to declare a variable in the shared memory.
- **`__device__`** to declare a variable in the global memory.
- **`__syncthreads()`** to sync the threads in the block.
- **`__threadfence()`** ensure that a writing operation is concluded before subsequent reads or writes from other threads can occur.

Atomic Operations

- Prevent race condition
- Can be performed on shared and global memory
- Hardware acceleration
- No synchronization needed
- Performance trade-off

Atomic operations in the code

CUDA provides many atomic operations:

- ***T atomicAdd(T* address, T val)***(or ***atomicSubs***): adds a value from a memory location and returns the old value.
- ***T atomicExch(T* address, T val)***: replaces the value at a memory location with a new value, and returns the old value.
- ***T atomicCAS(T* address, T compare, T val)***: compares the value at a memory location with an "expected" value and, if they are equal, swaps it with a new value. It returns the old value regardless of whether the swap occurred.

For the complete list of CUDA atomic operations: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

Example : Philosopher's dinner

```
const int NUM_PHILOSOPHERS = 5;

//Global variables
__device__ int d_forks[NUM_PHILOSOPHERS]; // Array of forks (0 = free, 1 = taken)
__device__ int d_philosophers[NUM_PHILOSOPHERS]; // Philosopher states (0 = thinking, 1 = eating)

//DEVICE code
__global__ void philosopher_dinner (int num_philosophers) { // Kernel function for philosopher
    behavior
    int id = threadIdx.x;

    // Each philosopher alternates between thinking and trying to eat
    while (true) {
        // Thinking
        d_philosophers[id] = 0; // Set philosopher state to thinking
        printf("Philosopher %d is thinking.\n", id);

        __syncthreads();
        __threadfence();

        // Try to pick up the two forks (left and right)
        int left_fork = id;
        int right_fork = (id + 1) % num_philosophers;

        // Pick up forks only if both are available (use atomic operations to prevent race
        conditions)
        if (atomicCAS(&d_forks[left_fork], 0, 1) == 0 && atomicCAS(&d_forks[right_fork], 0, 1) ==
0)
        {
            // Now the philosopher can eat
            d_philosophers[id] = 1; // Set philosopher state to eating
            printf("Philosopher %d is eating.\n", id);

            // Simulate eating
            for (int i = 0; i < 1000000; ++i); // Busy-waiting

            // Put down the forks after eating
            atomicExch(&d_forks[left_fork], 0);
            atomicExch(&d_forks[right_fork], 0);
        }
        else{
            atomicExch(&d_forks[left_fork], 0);
            atomicExch(&d_forks[right_fork], 0);
        }

        __syncthreads();
        __threadfence();
    }
}
```

```
//HOST code
int main() {
    // Initialize philosophers and forks arrays
    int h_forks[NUM_PHILOSOPHERS] = {0}; // 0 = fork is free
    int h_philosophers[NUM_PHILOSOPHERS] = {0}; // 0 = thinking, 1 = eating

    // Copy data to device
    cudaMemcpyToSymbol(d_forks, h_forks, sizeof(int) * NUM_PHILOSOPHERS);
    cudaMemcpyToSymbol(d_philosophers, h_philosophers,
        sizeof(int) * NUM_PHILOSOPHERS);

    // Launch kernel: each philosopher runs on its own CUDA thread
    philosopher_dinner <<<1, NUM_PHILOSOPHERS >>>(NUM_PHILOSOPHERS);

    // Wait for the kernel to complete
    cudaDeviceSynchronize();

    return 0;
}
```





Dynamic Parallelism

Dynamic Parallelism

Allow Kernel (*parent*) launch other kernels (*childs*)


- Runtime launch
- launched from the GPU
- High scalability
- Recursive algorithm optimisation
- Require explicit synchronization

Dynamic Parallelism in code


→ Simply call the child kernel inside the parent using `<<<>>>`

```
global void parentKernel() {  
    childKernel<<<blocks,threads>>>();  
}
```

CUDA programming model



CUDA dynamic parallelism



Streams and Event

CUDA streams introduce concurrency between different operations (kernels, memory transfers) on the GPU:

- kernels can be assigned to a specific stream
- Asynchronous execution
- Concurrency between streams

Synchronization can be achieved in different ways:

Explicit:

- Synchronize everything
- Synchronize w.r.t. a specific stream
- Using events

Implicit:

- Device memory allocation
- Non-Async version of memory operations
- Page-locked memory allocation

Streams in code

Streams

- **`cudaStreamCreate(&stream)/cudaStreamDestroy(stream)`**: to manage stream creation and destruction
- **`cudaStreamSynchronize(&stream)`**: blocks host until all CUDA calls in the stream are complete

Events

- **`cudaEventCreate(&event)/cudaEventDestroy(event)`**: to manage event creation and destruction
- **`cudaEventRecord(event, stream)`**: to signal completion of the event task
- **`cudaWaitEvent(stream2, event)`**: to wait an event before launch a new stream
- **`cudaEventSynchronize(event)`**: to synchronize computation till a specific event occurs

Pinned Memory

- **`cudaMallocHost(&host, size)`**: lock host memory to prevent paging.
- **`cudaMemCpyAsync()`**: copy values from host to device(and viceversa) asynchronously.

```
kernel1<<<,,,a>>>();
```

```
cudaEventRecord(e1, a);
```

```
kernel2<<<,,,b>>>();
```

```
cudaStreamWaitEvent(b, e1);
```

```
cudaMemcpyAsync(,,,b);
```

```
cudaEventRecord(e2, b);
```

```
kernel3<<<,,,a>>>();
```

```
cudaEventRecord(e3, a);
```


```
cudaEventSynchronize(e2);
```

CPU code



```
cudaStreamWaitEvent(b, e3);
```

```
kernel4<<<,,,b>>>();
```

2 GPU streams



1 CPU Threads





Final Consideration

Final Considerations

- *Memory management* to minimize transfer between CPU(*host*) and GPU(*device*).
- *Thread hierarchy design* is crucial, carefully design grid and block dimensions.
- Correct usage of *Atomic operations* and *synchronization* to avoid race conditions and maintain parallelism.
- Maximize parallelism using *dynamic parallelism* and *streams*.
- Performance profiling: NVIDIA provides tools like *Nsight* to identify bottlenecks and optimize performance.



END

Thanks for your attention

References

- Fang, Jianbin, et al. "Parallel programming models for heterogeneous many-cores: a comprehensive survey." *CCF Transactions on High Performance Computing* 2 (2020): 382-400.
- <https://ams148-spring18-01.courses.soe.ucsc.edu/home.html> (Lectures 2, 3 and 8)
- Nickolls, John, et al. "Scalable parallel programming with cuda: Is cuda the parallel programming model that application developers have been waiting for?." *Queue* 6.2 (2008): 40-53.
- <https://developer.nvidia.com/blog/cuda-dynamic-parallelism-api-principles/>